

Sistemas Distribuídos

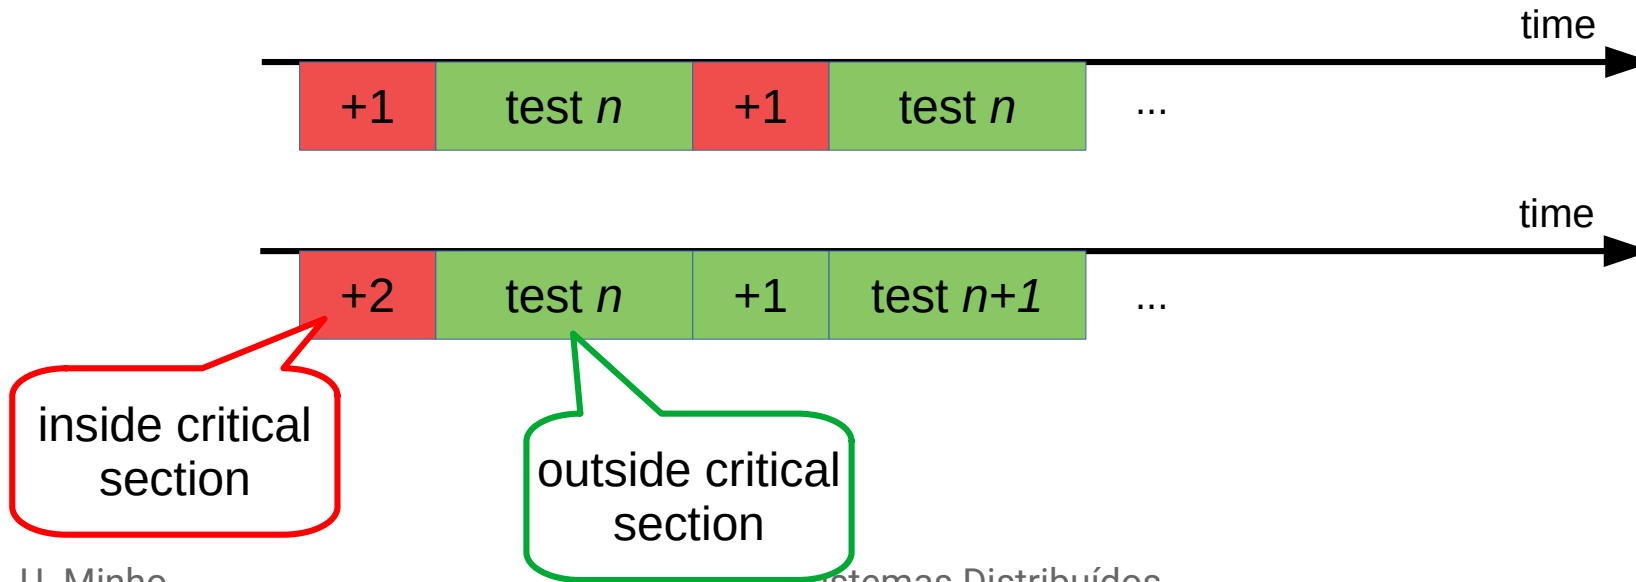
José Orlando Pereira

Departamento de Informática
Universidade do Minho




Motivation

- Consider two versions of the parallel primality testing code:
 - Increment $+1$ and get n , test n
 - Increment $+2$ and get n , test n and $n+1$




Motivation



- Eventually, at least one thread is blocked waiting for mutex...

Motivation



- Reducing the contention on critical sections lessens the performance impact of synchronization

Roadmap

- Use synchronization primitives to write correct concurrent code and avoid busy waiting
- Need to minimize **time** in critical sections
- Need to minimize **contention** in critical sections

Example: Game



Game state and operations

- State:
 - `Map<String,Player> players;`
 - `class Player {`
 - `int x,y;`
 - `int life, score;``}`
- Operations:
 - drop in the game, move, and shoot
 - draw the game

First approach

- 1 thread for each player^(*)
- 1 lock for the shared game state

^(*) Later we make it distributed...

First approach

- ```
void draw() {
 try { l.lock();
 players.values().forEach(p→Draw3D(p.x, p.y));
 } finally { l.unlock(); }
}
```

Try/finally make it  
work with exceptions

Lengthy computation  
inside critical section

- Problems:
  - Either drawing or moving
  - Drawing takes a long time
  - “Lag”...

# Immutable objects

- class Coord { final int x, y; }
- class Player {  
    Coord xy;  
    int life, score;  
}
- void draw() {  
    { try { l.lock();  
        c=players.values().stream()  
          .map(p→p.xy).collect(toList());  
        } finally { l.unlock(); }  
    c.forEach(c→Draw3D(c.x, c.y));  
}

All fields final

Lengthy computation  
outside critical section

# Multiple locks

- Can't move two players concurrently
- Forget “drop in the game” for now...
- Use one lock for each player:
- ```
class Player {  
    Lock l;  
    Coord xy;  
    int life, score;  
}
```

Multiple locks


- ```
void move(...) {
 try { l.lock();
 xy = new Coord(...);
 } finally { l.unlock(); }
}
```
- ```
Coord getLocation() {  
    try { l.lock();  
        return xy;  
    } finally { l.unlock(); }  
}
```

Multiple locks

- ```
void shoot(String sn, String tn) {
 Player s = players.get(sn);
 Player t = players.get(tn);
 try { s.l.lock(); t.l.lock();
 t.life--;
 s.score++;
 } finally { t.l.unlock(); s.l.unlock(); }
}
```

# Deadlock

- What if two players shoot at each other simultaneously ( $A \rightarrow B$  and  $B \rightarrow A$ ) ?
- What if  $A \rightarrow B$ ,  $B \rightarrow C$  and  $C \rightarrow A$ ?
- What if ...



# Lock ordering

- What if two players A, B shoot at each other simultaneously?
  - A acquires A, B
  - B acquires A, B
- What if  $A \rightarrow B$ ,  $B \rightarrow C$  and  $C \rightarrow A$ ?
  - A acquires A, B
  - B acquires B, C
  - C acquires A, C

# Lock ordering

- ```
void shoot(String sn, String tn) {  
    Player s = players.get(sn);  
    Player t = players.get(tn);  
    try { Stream.of(sn,tn).sorted()  
        .forEach(n→players.get(n).l.lock());  
        t.life--;  
        s.score++;  
    } finally { t.l.unlock(); s.l.unlock(); }  
}
```

Acquire locks
in a fixed order

Release in
any order

Fairness

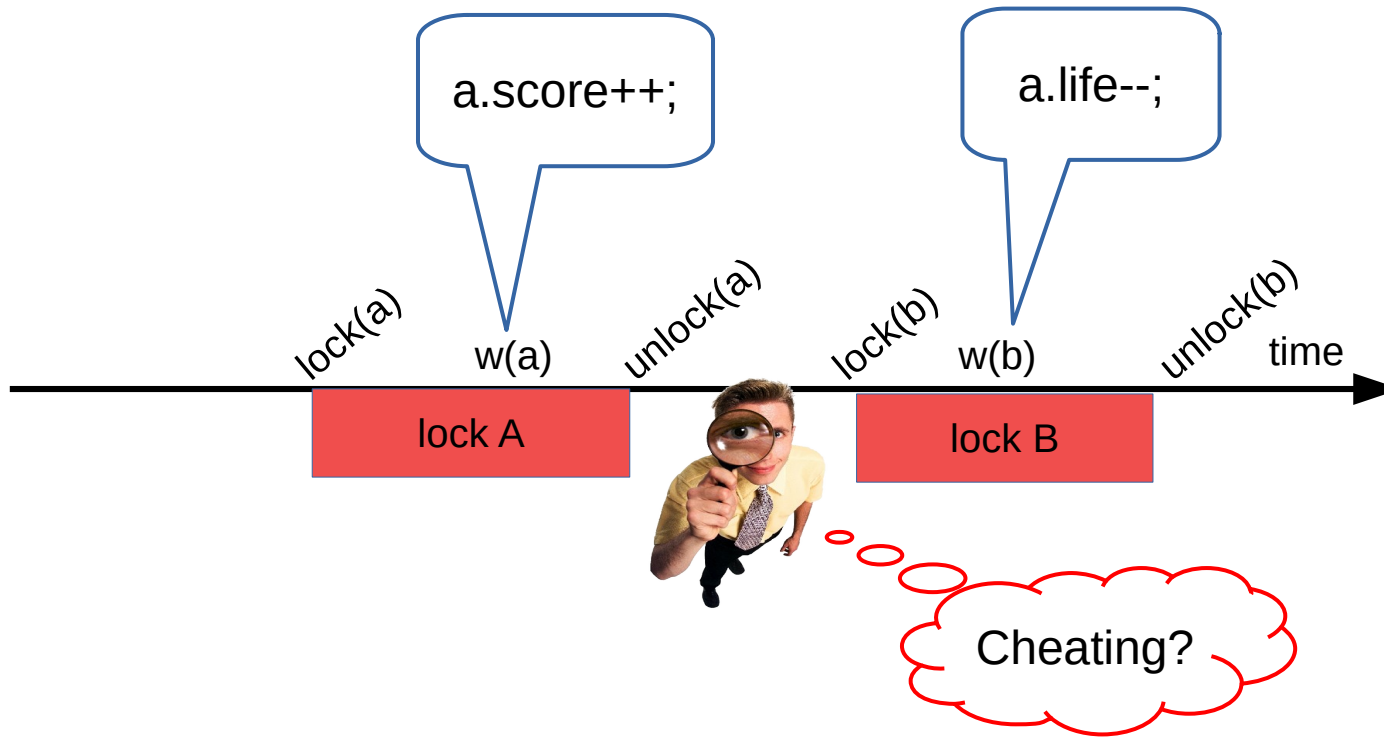
- “Doesn’t lock ordering mean that player A has an advantage?”
- No. It means that:
 - When A shoots some X and X shoots A, at the same time, the winner will be decided by lock of A
 - Threads acquiring the same lock are optionally fair with j.u.c. ReentrantLock (i.e. approximately FIFO)
- So they have the same chances regardless of the lock used

Multiple locks

- Acquiring all locks needed at the start and releasing them at the end of an operation works as well as single global lock
- What if we need to read some data before acquiring further locks?
- How to further reduce the time holding locks?


Multiple locks

- Why acquire both locks simultaneously?
 - If we don't....



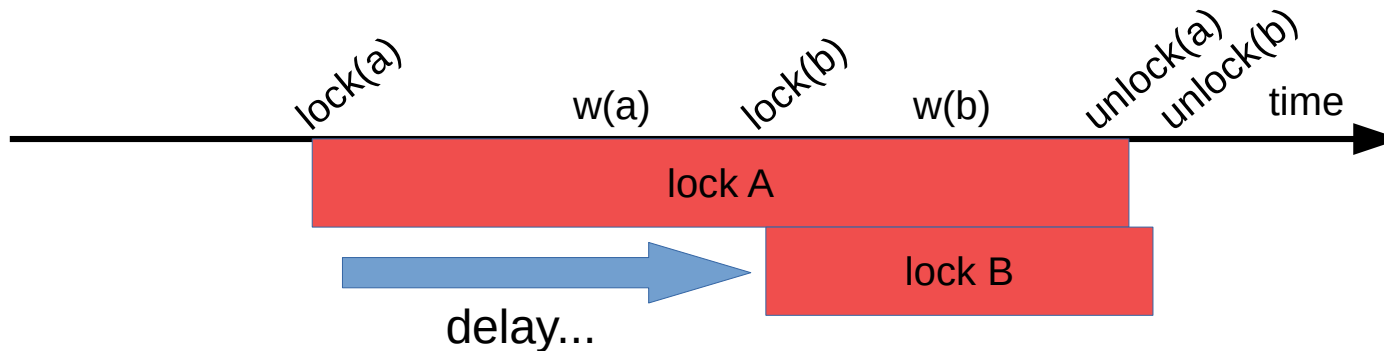
Multiple locks

- Why acquire both locks simultaneously?
(The observer will also lock A and B.)



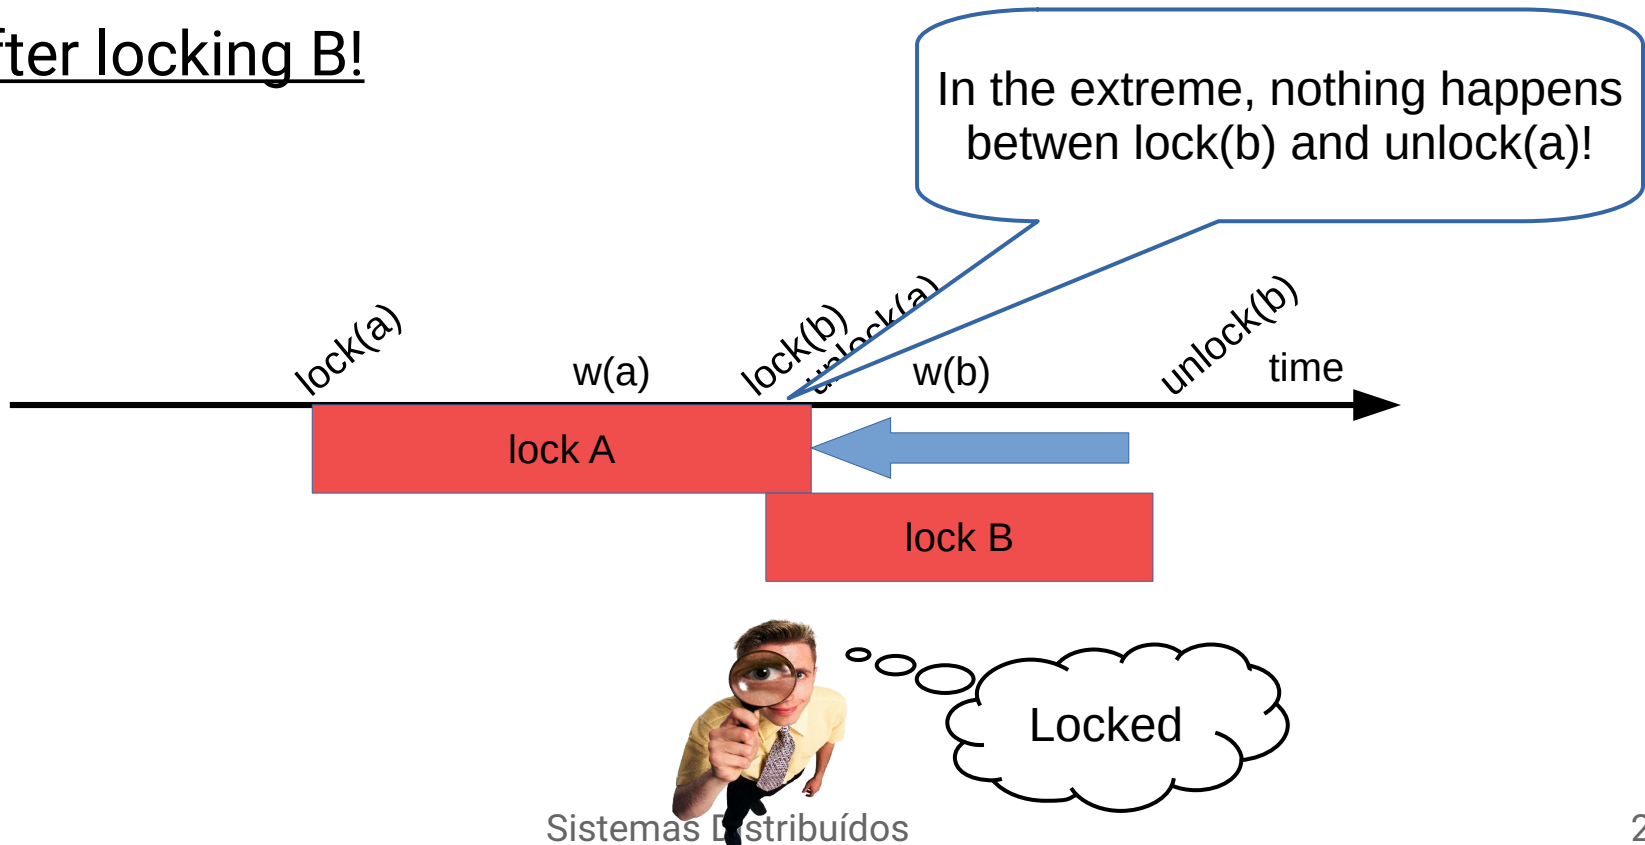
Lock later

- How much can we delay acquiring lock for B?
 - Until needed for modifying item b




Unlock earlier

- How much can we anticipate releasing lock for A?
 - After modifying item a and...
 - after locking B!



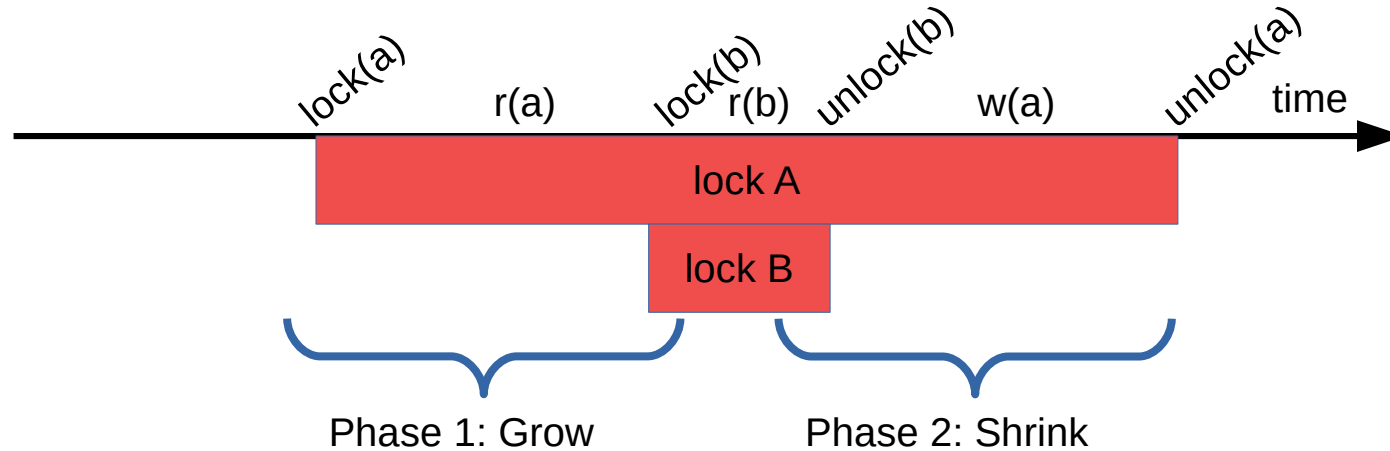
Two phase locking (2PL)

- **Rule 1:** All lock() precede all unlock()
- **Rule 2:** Each data item is read/written within the corresponding lock
 - Equivalent to holding all relevant locks, all the time



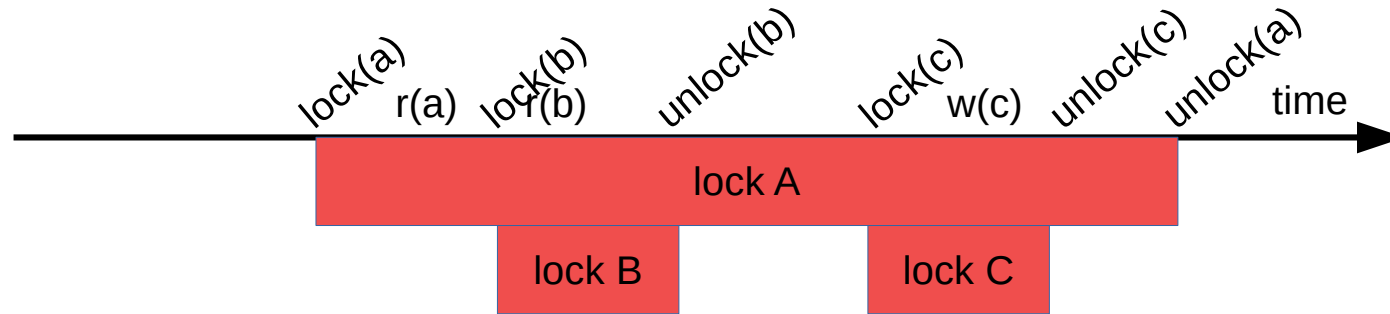
Two phase locking (2PL)

- Another example:

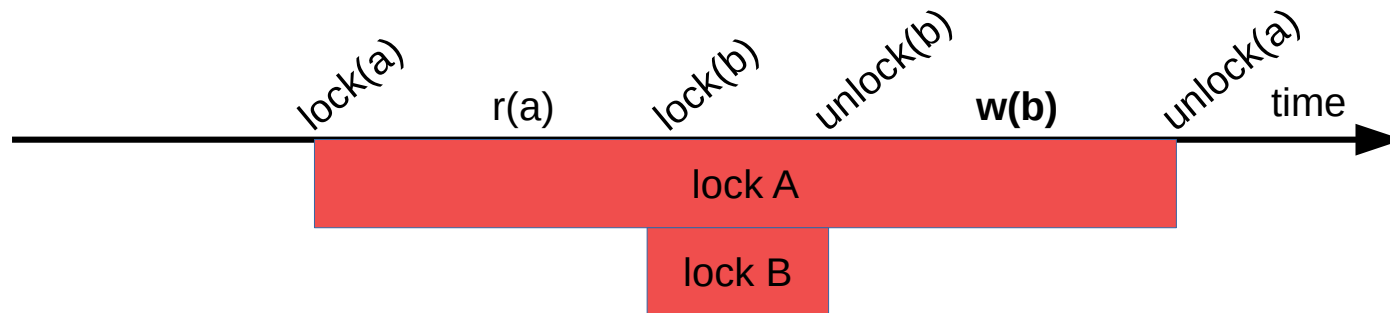


Two phase locking (2PL)

- Fails Rule 1:



- Fails Rule 2:



Two phase locking (2PL)

- ```
void shoot(String sn, String tn) {
 Player s = players.get(sn);
 Player t = players.get(tn);
 Stream.of(sn,tn).sorted()
 .forEach(n→players.get(n).l.lock());
 t.life--;
 t.l.unlock();
 s.score++;
 s.l.unlock();
}
```
- Diagram illustrating the Two Phase Locking (2PL) protocol for the `shoot` method:
- Phase 1: Grow** (Expansion Phase):
    - Acquiring locks on `sn` and `tn`: `Stream.of(sn,tn).sorted().forEach(n→players.get(n).l.lock());`
  - Phase 2: Shrink** (Contraction Phase):
    - Releasing locks: `t.l.unlock();`, `s.score++;`, and `s.l.unlock();`
- Note: The lock acquisition and release operations are marked with blue arrows indicating the flow of the protocol.

# Collection locking

- What if the collection is not immutable?
  - “drop in the game”
- Add back a global lock to game state...

# Collection locking

- ```
void shoot(String sn, String tn) {  
    l.lock();  
    Player s = players.get(sn);  
    Player t = players.get(tn);  
    Stream.of(sn,tn).sorted()  
        .forEach(n→players.get(n).l.lock());  
    t.life--;  
    s.score++;  
    t.l.unlock(); s.l.unlock();  
    l.unlock();  
}
```

Collections with 2PL

- void shoot(String sn, String tn) {
 l.lock();
 Player s = players.get(sn);
 Player t = players.get(tn);
 Stream.of(sn,tn).sorted()
 .forEach(n→players.get(n).l.lock(););
 l.unlock();
 t.life--;
 t.l.unlock();
 s.score++;
 s.l.unlock();
}

Is ordering needed?

Phase 1: Grow

Phase 2: Shrink

Collections with 2PL

- void shoot(String sn, String tn) {
 l.lock();
 Player s = players.get(sn);
 Player t = players.get(tn);
 s.l.lock();
 p.l.lock();
 l.unlock();
 t.life--;
 t.l.unlock();
 s.score++;
 s.l.unlock();
}

No, if these locks are always acquired in the context of the collection lock!

Conclusions

- Minimizing critical sections is key to performance and scale
- Strategies to reduce impact of critical sections:
 - Immutable objects
 - Granular locking
 - Two phase locking
 - Collections
- Avoid deadlocks by using a fixed locking order

Locks vs Variables

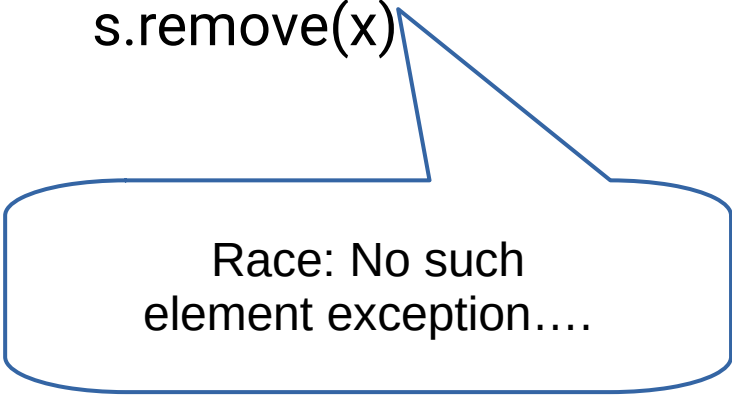
- “Which lock corresponds to each data item?”
- Multiple threads accessing some data item concurrently must have acquired the same lock
- Not automatic / not checked
- It is up to the developer to ensure this!

Pitfall: Encapsulated locks

- Keep variables and the corresponding lock encapsulated within the same object
- (The default using old-style “synchronized” in Java.)

Wrong

```
class SomeClass {  
    SomeState s;  
    void doSomething() {  
        if (s.contains(x))  
            s.remove(x)  
    }  
}
```



Race: No such
element exception....

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

Solution

```
class SomeClass {  
    private Lock l;  
    SomeState s;  
    void doSomething() {  
        l.lock();  
        if (s.contains(x))  
            s.remove(x)  
        l.unlock();  
    }  
}
```

Now useless...

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

Better solution

```
class SomeClass {  
    private Lock l;  
    SomeState s;  
    void doSomething() {  
        l.lock();  
        if (s.contains(x))  
            s.remove(x)  
        l.unlock();  
    }  
}
```

```
class SomeState {  
    private Lock l;  
    boolean contains(...) {  
        l.lock(); ... l.unlock();  
    }  
    void remove(...) {  
        l.lock(); ... l.unlock();  
    }  
}
```

Rely on locking by the callers. This is done by Java Collections (Lists, ...)


Pitfall: Shared vs thread-local state

- Program state often contains:
 - Local thread state in workers
 - Shared state, used by all threads
- Both are objects, with instance variables

Wrong

```
class Worker
  extends Thread {
    Lock l;
    SharedState s;
    void doSomething() {
      l.lock(); s.doit(); l.unlock();
    }
    public void run() {
      doSomething();
    }
  }
```


```
class SharedState {
  public void doit() {
    ...
  }
}
```



Solution

```
class Worker
  extends Thread {
    SharedState s;
    void doit() {
      s.doit();
    }
    public void run() {
      doit();
    }
  }
}
```

```
class SharedState {
  Lock l;
  public void doit() {
    l.lock(); ... l.unlock();
  }
}
```



Pitfall: Class/global variables

- Variables marked with “static” in Java are global and (probably) need concurrency control
 - Not if marked “final”
 - Not if the class is used by a single thread

Wrong

```
class SomeClass {  
    private static int s;  
    void doSomething() {  
        s = s+1;  
    }  
}
```



Race!

Still wrong

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private static int s;  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```

There is one lock for each object, but s is shared!

Solution

```
class SomeClass {  
    private static Lock l = new ReentrantLock();  
    private static int s;  
    void doSomething() {  
        l.lock();  
        s = s+1;  
        l.unlock();  
    }  
}
```

Collections

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    List getElements() {  
        try { l.lock();  
            return l;  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
List l = l.getElements();  
l.add(...);
```



Race!

Collections

```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    Iterator getElements() {  
        try { l.lock();  
            return l.iterator();  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
    Iterator i = l.getElements();  
    while(i.hasNext())
```

...



Race!

Collections


```
class SomeClass {  
    private Lock l = new ReentrantLock();  
    private List l;  
    List getElements() {  
        try { l.lock();  
            return l.clone();  
        } finally { l.unlock(); }  
    }  
}
```

```
SomeClass s = ...;  
List l = l.getElements();  
l.add(...);
```

Not adding to
the list...

Summary

- There is no simple rule to match locks with variables
- Some thinking needed... :-)



Scaling up

- Example:
 - In a distributed database table with millions of records
 - Executing “select sum(x) from ... where ...” queries
 - Updating records
- Do we use a single lock?
 - Cannot run more than one query at the same time
- Do we use a lock for each line?
 - Way too many individual locks!

Readers-Writers locks

- Strict mutual exclusion with locks is too conservative:
 - More than one reader would not be a problem
 - A writer must exclude all others (readers and writers)
- Different methods for readers and writers:

```
interface ReadWriteLock {  
    Lock readLock();  
    Lock writeLock();  
}
```



Readers-Writers fairness

- Give preference to readers
 - Allow more readers in, even if a writer is waiting
 - The writer may starve...
- Give preference to writers
 - Do not allow more readers in if a writer is waiting
 - Less concurrency among readers
- Fair and efficient:
 - Readers and writers in FIFO order
 - Allow readers to skip up to k writers in the queue

Revisiting collections with 2PL

- ```
void shoot(String sn, String tn) {
 l.readLock().lock();
 Player s = players.get(sn);
 Player t = players.get(tn);
 Stream.of(sn,tn).sorted()
 .forEach(n→players.get(n).l.lock());
 l.readLock().unlock();
 t.life--;
 t.l.unlock();
 s.score++;
 s.l.unlock();
}
```

Allow multiple  
threads to acquire  
locks concurrently

Sorting is needed  
again

# Lock managers

- Individual locks inefficient for huge collections of objects
  - A lock object uses memory even when not in use
- A lock manager provides locks on demand:

```
interface LockManager {
 void lock(Object name);
 void unlock(Object name);
}
```

lookup lock l for "name" in map  
if it doesn't exist:  
    create it and add to map  
l.lock()!

lookup lock for "name"  
l.unlock()  
if nobody else is using it:  
    remove it from map

# Lock managers

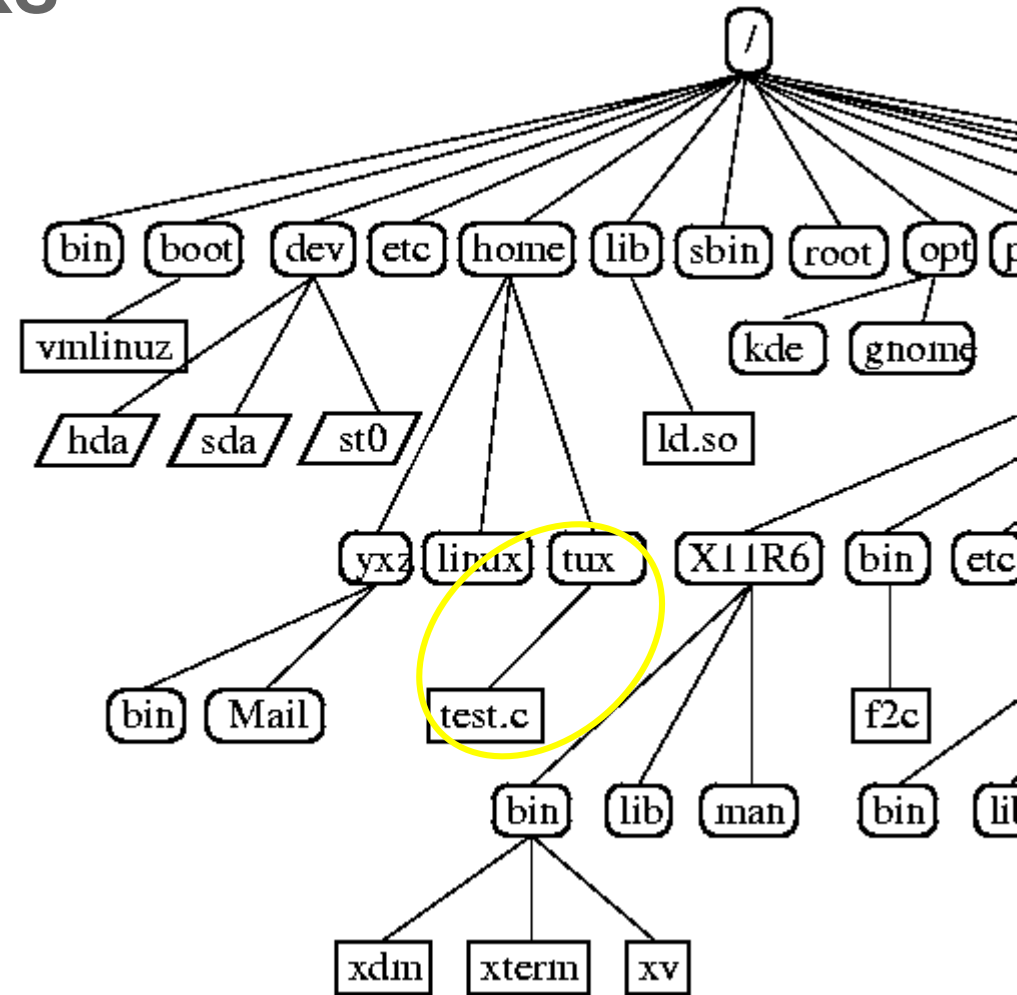
- Usually provides Readers-Writers semantics:  
(Reader = Shared, Writer = Exclusive)

```
enum Mode { SHARED, EXCLUSIVE };
```

```
interface LockManager {
 void lock(Object name, Mode mode);
 void unlock(Object name);
}
```

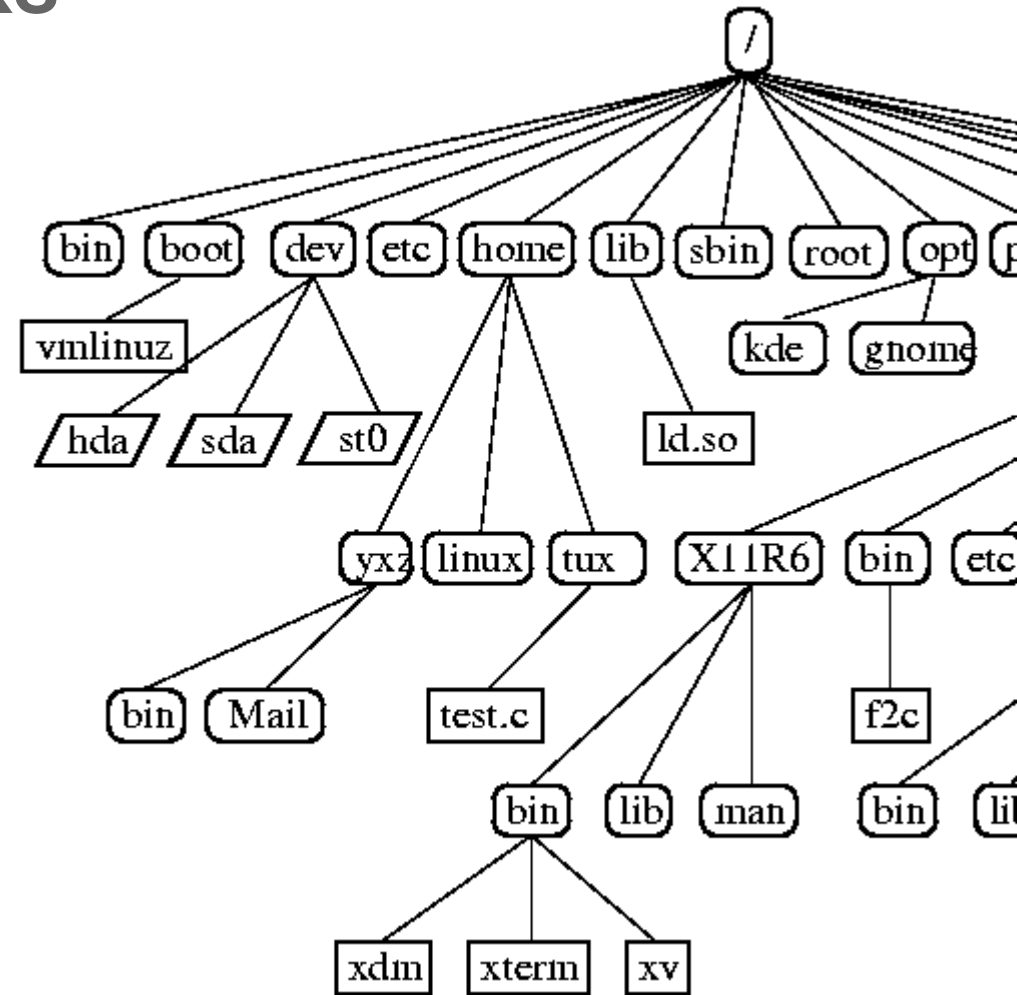
# Multiple granularity locks

- Motivation:
  - Locking `/home/tux/*`
  - Assume large number of files
- Inefficient even with a lock manager
- Idea: Take advantage of hierarchical namespace and lock folders



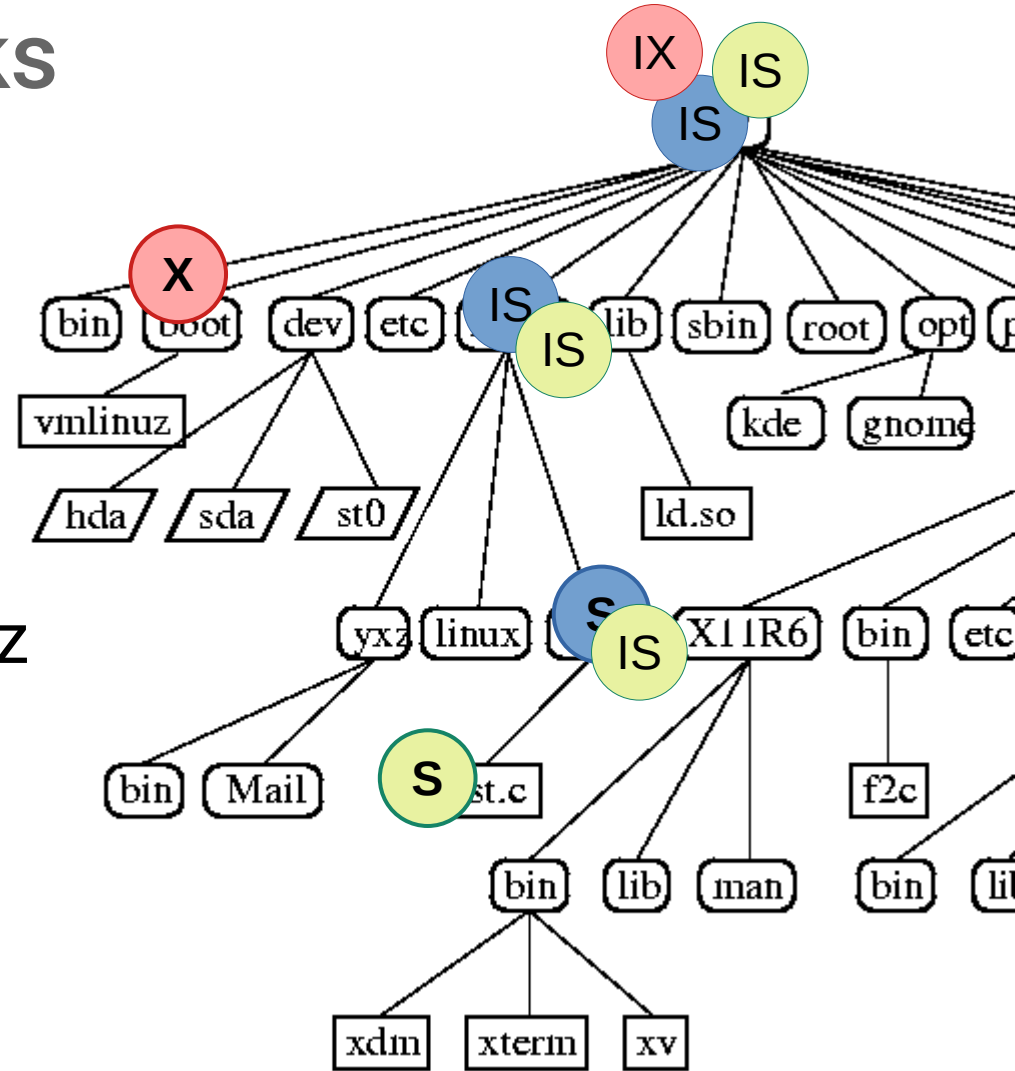
# Multiple granularity locks

- Protocol:
  - “intention” locks on containers
  - “actual” locks on the target
- Intention locks conflict with actual locks, not with other intention locks
- Combine with (S)hared and e(X)clusive semantics



# Multiple granularity locks

- Shared lock /home/tux
- Shared lock on /home/tux/test.c
- Exclusive lock on /boot
  - Shared lock on /boot/vmlinuz
    - IS on /boot/ conflicts with X
  - Exclusive lock on /home
    - X on /home conflicts with IS





# Compatibility matrix

- An MGL is defined by a compatibility matrix:

The diagram shows a 5x5 compatibility matrix. A blue bracket on the left, labeled 'acquired', spans the rows. A blue bracket on top, labeled 'requesting...', spans the columns. The matrix cells are colored: yellow for 'Yes' and red for 'No'.

| Mode | IS  | IX  | S   | X  |
|------|-----|-----|-----|----|
| IS   | Yes | Yes | Yes | No |
| IX   | Yes | Yes | No  | No |
| S    | Yes | No  | Yes | No |
| X    | No  | No  | No  | No |