

Sistemas Distribuídos

José Orlando Pereira

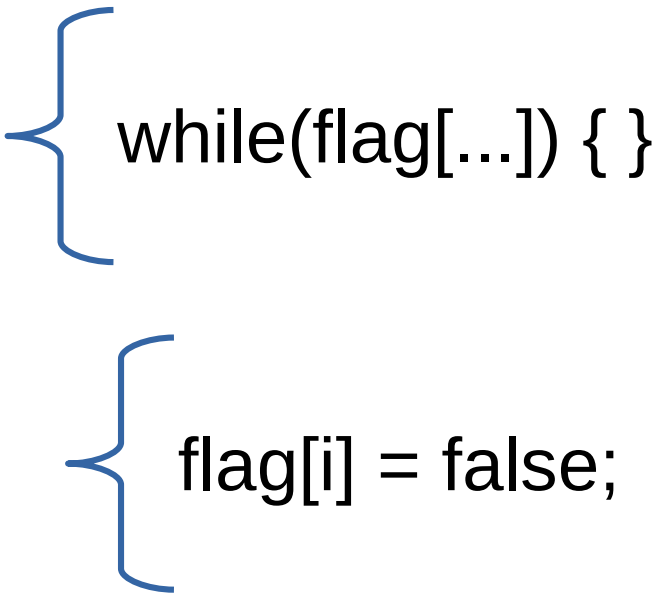
Departamento de Informática
Universidade do Minho



Mutex with Peterson's/...

- Does it really work?

```
try {  
    l.lock();  
    c=c+1;  
} finally {  
    l.unlock();  
}
```



```
while(flag[...]) { }
```

```
flag[i] = false;
```

Quiz

- Two variables:
 - `int i=0, j=0;`
- Writer code:
 - `i=1; j=1;`
- Reader code:
 - `rj=j; ri=i; System.out.println(rj+", "+ri);`
- Possible results:
 - a) 0, 0 ✓
 - b) 1, 1 ✓
 - c) 0, 1 ✓
 - d) 1, 0 ✓




Why!?!?

} running
concurrently!

Memory order

- Steps to write a variable:
 1. Write to cache
 2. Flush cache to memory



(Image from <http://booksite.elsevier.com/9780123705914/?ISBN=9780123705914> . CC By-SA-3.0.)

Memory order

- Possible outcome with two variables:


1. Write i to cache

2. Write j to cache

3. Flush j from cache to memory

5. Flush i from cache to memory


4. Paradox observed if i,j read here!!




(Image from <http://booksite.elsevier.com/9780123705914/?ISBN=9780123705914> . CC By-SA-3.0.)

Consequence

```
try {  
    l.lock();  
    c=c+1;  
} finally {  
    l.unlock();  
}
```

 while(flag[...]) { }


 flag[i] = false;

- Initially c=10
- One thread:
 - read c
 - write c = 11
 - write flag[i] = false
- Other thread:
 - read flag[...] = false
 - read c = 10!!!!



Solution: Memory barriers

- Declare: volatile int j; or AtomicInteger j;
- Reading from a volatile j waits for all writes preceding the observed value on j to be also visible
 - Writer code:
 - i=1; j=1;
 - Reader code:
 - rj=j; ri=i; System.out.println(rj+", "+ri);



waits for write i=1
to be visible

Should I use it?

- Synchronization operations also act as memory barriers:
 - `java.util.concurrent.*` primitives
 - `synchronized` keyword
- Volatile variables also impact performance:
 - A volatile access takes as much time as using a lock
 - Therefore... **use the lock!**


Corollary

```
class X {  
    private Y y;  
    void changeY() {  
        l.lock();  
        tmp.i = 1;  
        y = tmp;  
        l.unlock();  
    }  
    int getY() {return y.i;}  
}
```

- Can we omit synchronization in getters?

Corollary

```
class X {  
    private Y y;  
    void changeY() {  
        l.lock();  
        tmp.i = 1;  
  
        y = tmp;  
        l.unlock();  
    }  
    int getI() { return y.i;}  
}
```



- Can we omit synchronization in getters?
 - **NO!**
- Can read inconsistent Y fields!
- In this case:
 - reader might not see `y.i == 1!!!!`

The real world

- Is j.u.c.ReentrantLock implemented with Bakery or something similar?
 - **No!**
 - Needs $O(n)$ space for each lock (with n threads)
 - Consumes CPU time / power when busy
- These algorithms are a good for:
 - Understanding concurrent programming and races
 - Preparing for distributed algorithms

The real world

- `j.u.c.ReentrantLock` is implemented with atomic operations and scheduling:
 - Atomic operations (`testAndSet`, `compareAndSwap`, ...) are used to quickly check if mutex is available
 - For more on atomic operations, see Chapter 7 of TAOMP (not part of the program!)
 - The thread is suspended by the OS when waiting for longer periods
 - Remember scheduling from Operating Systems
 - `man futex` on Linux

j.u.c Locks vs Monitors

```
class C {  
    private int i;
```

There is a hidden “lock” in each object used by “synchronized”

```
synchronized public void m() {
```

```
    i++;
```

```
}
```

```
}
```

```
class C {  
    private int i;  
    private Lock l =  
        new ReentrantLock();
```

```
    public void m() {  
        try { l.lock();  
            i++;  
        } finally { l.unlock(); }
```

```
}
```

```
}
```

Equivalent code
(approximately...)

j.u.c. Locks vs Monitors

- Main differences, for now:
 - Synchronized blocks are nested in LIFO order
vs.
j.u.c. Locks can be unlocked in any order
 - To take advantage of two phase locking
 - Threads waiting for a synchronized block enter in any order
vs.
j.u.c. Lock can be configured for threads waiting to enter in a fair (not FIFO!) order
- More later...

j.u.c Conditions vs Monitors

- class C {

There is a hidden “condition” in each object used by “wait()/notifyAll()”

```
synchronized public void m1() {  
    while(...) wait();  
}
```

```
synchronized public void m2() {  
    notifyAll();  
}
```

- class C {

```
private Lock l =  
    new ReentrantLock();  
private Condition c =  
    l.newCondition();  
public void m1() {  
    try { l.lock();  
        while(...) c.await();  
    } finally { l.unlock(); }  
}  
public void m2() {  
    try { l.lock();  
        c.signalAll();  
    } finally { l.unlock(); }  
}
```

Equivalent code
(aproximately...)

j.u.c. Conditions vs Monitors

- Main differences:
 - One implicit condition for each lock
vs.
Many j.u.c. conditions for the same lock
 - Avoids signalAll()
 - Threads waiting for a condition wakeup in any order
vs.
Threads waiting for a j.u.c. Condition obtained from a ReentrantLock wakeup in FIFO order (but may not acquire lock in FIFO order...)